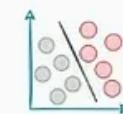


Pandas ↔ Polars ↔ SQL ↔ PySpark



blog.DailyDoseofDS.com

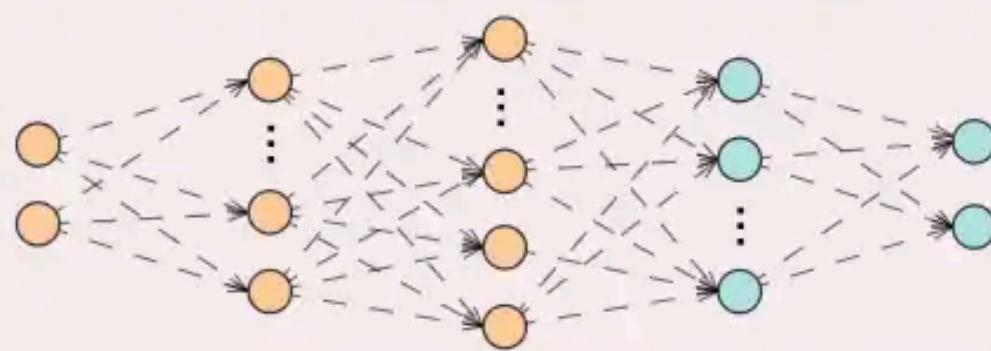
Operation	Pandas	Polars	SQL	PySpark
Import	<code>import pandas as pd</code>	<code>import polars as pl</code>	-	<code>from pyspark.sql import SparkSession spark = SparkSession.builder.appName("ABCD")</code>
Read CSV	<code>df = pd.read_csv(file)</code>	<code>df = pl.read_csv(file)</code>	<code>LOAD DATA INFILE 'data.csv' INTO TABLE table FIELDS TERMINATED BY ',' LINES TERMINATED BY '\n' IGNORE 1 ROWS;</code>	<code>df = spark.read .csv("data.csv")</code>
Print first 10 (or k) rows	<code>df.head(10)</code>	<code>df.head(10)</code>	<code>SELECT * FROM table LIMIT 10;</code>	<code>df.show(10)</code>
Dimensions	<code>df.shape</code>	<code>df.shape</code>	<code>SELECT count(*) FROM table;</code> <code>SELECT count(*) FROM INFORMATION_SCHEMA.COLUMNS where TABLE_NAME = 'table';</code>	<code>df.count()</code> <code>len(df.columns)</code>
Datatype	<code>df.dtypes</code>	<code>df.dtypes</code>	<code>DESCRIBE table;</code>	<code>df.printSchema()</code>
Select column(s)	<code>df[["col1", "col2"]]</code>	<code>df[["col1", "col2"]]</code>	<code>SELECT column FROM table;</code>	<code>df.select("col1", "col2")</code>
Filter Data	<code>df[df.column > 10]</code>	<code>df[df.column > 10]</code> <code>df.filter(pl.col("column") > 10)</code>	<code>SELECT * FROM table where column>10;</code>	<code>df.filter(df["column"]>10)</code>
Sort	<code>df.sort_values("column")</code>	<code>df.sort("column")</code>	<code>SELECT * FROM table ORDER BY column;</code>	<code>df.orderBy("column")</code>
Fill NaN	<code>df.column.fillna(0)</code>	<code>df.column.fill_nan(0)</code>	<code>UPDATE table SET column=0 WHERE column IS NULL;</code>	<code>df.na.fill(0)</code>
Join	<code>pd.merge(df1, df2, on ="col", how="inner")</code>	<code>df1.join(df2, on="col", how="inner")</code>	<code>SELECT * FROM table1 JOIN table2 ON (table1.col = table2.col);</code>	<code>df1.join(df2, on="col", how="inner")</code>
Concatenate	<code>pd.concat((df1, df2))</code>	<code>pl.concat((df1, df2))</code>	<code>SELECT * FROM table1 UNION ALL table2;</code>	<code>df1.union(df2)</code>
Group	<code>df.groupby("column"). agg_col.mean()</code>	<code>df.groupby("column"). agg(pl.mean("agg_col"))</code>	<code>SELECT column, avg(agg_col) FROM table GROUP BY column;</code>	<code>df.groupBy("column"). agg(avg("agg_col"))</code>
Unique values	<code>df.column.unique()</code>	<code>df.column.unique()</code>	<code>SELECT DISTINCT column FROM table;</code>	<code>df.select("column"). distinct()</code>
Rename column	<code>df.rename(columns = {"old_name": "new_name"})</code>	<code>df.rename(mapping = {"old_name": "new_name"})</code>	<code>ALTER TABLE table RENAME COLUMN old_name TO new_name;</code>	<code>df.withColumnsRenamed({"old_name": "new_name"})</code>
Delete column	<code>df.drop(columns = ["column"])</code>	<code>df.drop(name = ["column"])</code>	<code>ALTER TABLE table DROP COLUMN column;</code>	<code>df.drop("col1", "col2")</code>

Next... 

4 Strategies for Multi-GPU Training

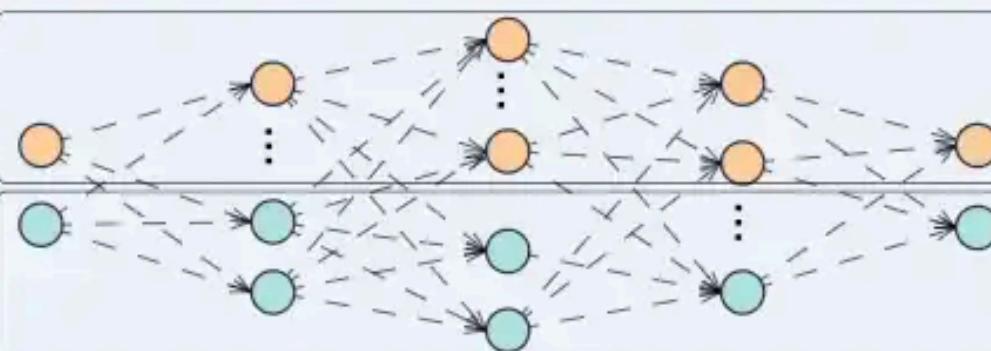
blog.DailyDoseofDS.com

Model parallelism



- Layer on 1st GPU
- Layer on 2nd GPU

Tensor parallelism

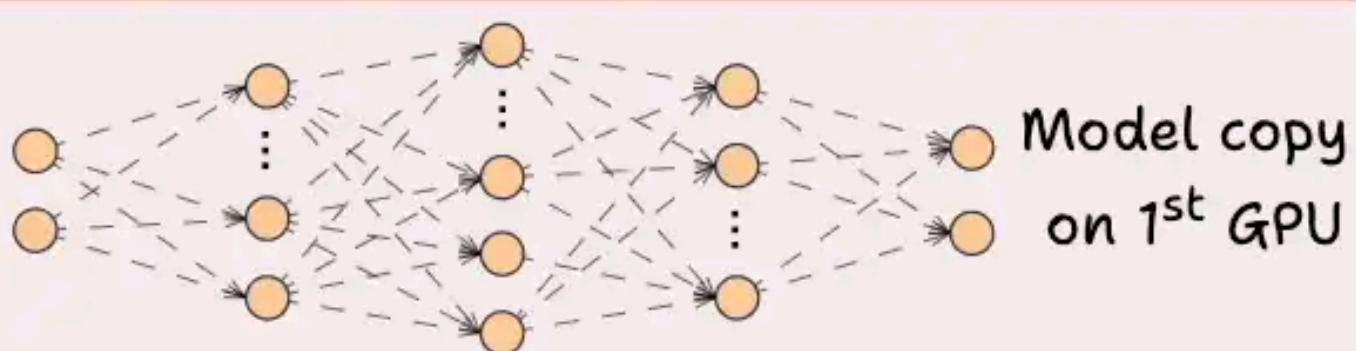


- Neurons on 1st GPU
- Neurons on 2nd GPU

Data parallelism

Data subset #1

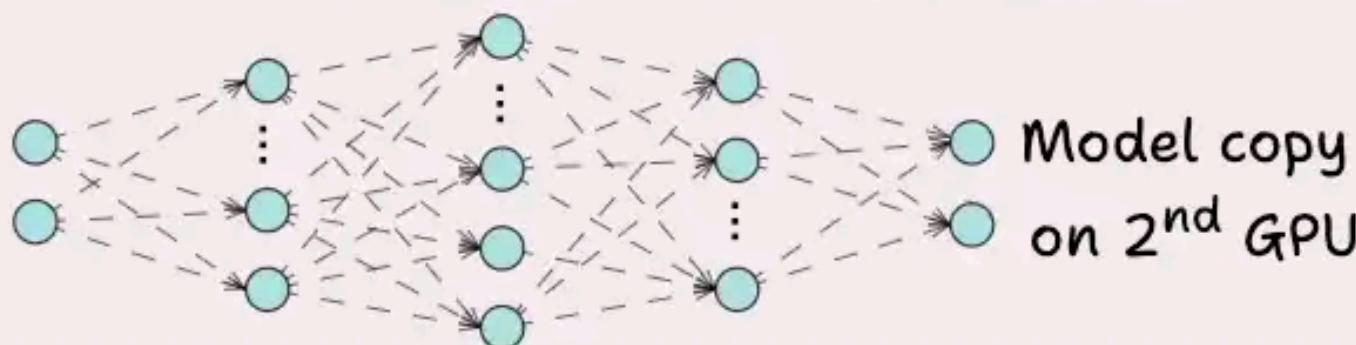
x1	x2	x3
...
...
...



- Model copy
- on 1st GPU

Data subset #2

x1	x2	x3
...
...



- Model copy
- on 2nd GPU

Pipeline parallelism

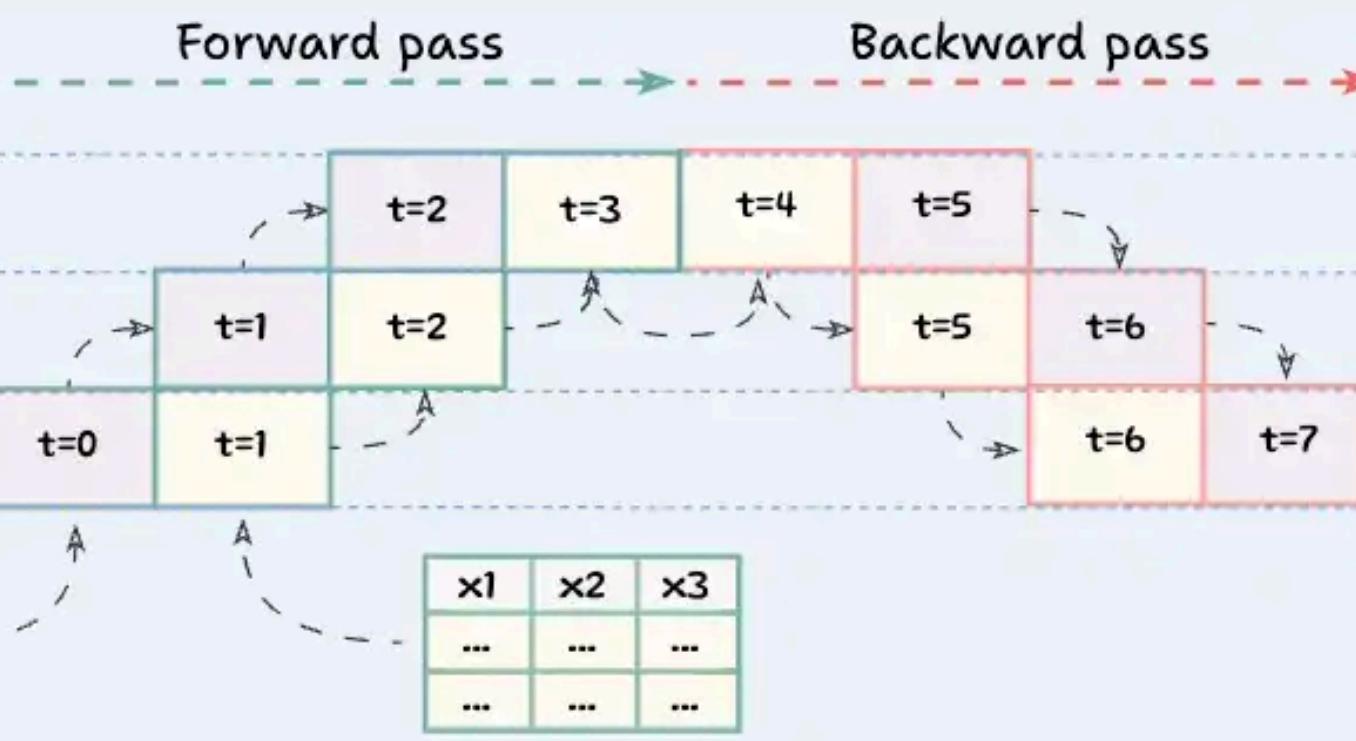
Forward pass

Backward pass

3rd layer
on 3rd GPU
2nd layer
on 2nd GPU
1st layer
on 1st GPU

x1	x2	x3
...
...

Data subset #1

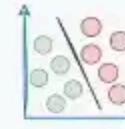


x1	x2	x3
...
...

Data subset #2

Next...

4 Ways to Test ML Models in Production

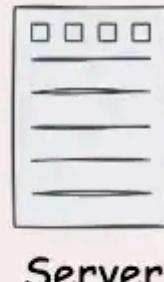


blog.DailyDoseofDS.com

1) A/B Testing



Request
--->



Server

90%
--->
10%
--->



Legacy model



Candidate model

Send some requests to the candidate model.

2) Canary Testing



Server



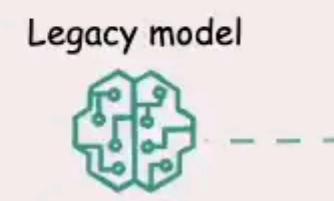
Legacy model



Candidate model

Release the candidate model only to a few users.

3) Interleaved Testing

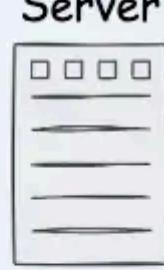


Use both models for predictions.

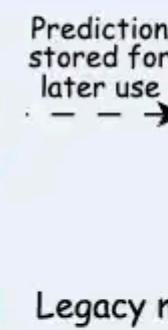
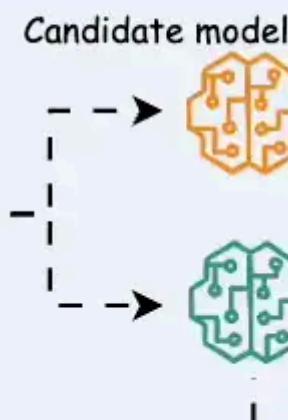
4) Shadow Testing



Request
--->



Server



Prediction stored for later use



Store the predictions from candidate model for later use/inspection.

Next...

15 Ways to Optimize Neural Network Training



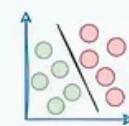
blog.DailyDoseofDS.com

- Use efficient optimizers—AdamW, Adam, etc.
- Utilize hardware accelerators (GPUs/TPUs).
- Max out the batch size.
- Use Bayesian Optimization if hyperparameter search space is big.
- Set max_workers in the Dataloader.
- Set pin_memory in Dataloader.
- Use mixed precision training.
- Use He or Xavier initialization for faster convergence (usually helps).
- Use activation checkpointing to optimize memory (run-time will go up).
- Utilize multi-GPU training through Model/Data/Pipeline/Tensor parallelism.
- For large models, use DeepSpeed, FSDP, YaFSDP, etc.
- Normalize data after transferring to GPU (for numerical data, like pixels).
- Use gradient accumulation (may have marginal improvement at times).
- Always use DistributedDataParallel, not DataParallel.
- `torch.rand(2, 2, device = ...)` creates tensors directly on the GPU.
- ~~`torch.rand(2,2).cuda()`~~ first creates on the CPU, then transfers to GPU.

→ Always profile your code to identify and to identify eliminate performance bottlenecks. ←

→ Next...

Time Complexity of 10 Most Popular ML Algorithms



blog.DailyDoseofDS.com

Training

Inference

	Linear Regression (OLS)	$O(nm^2 + m^3)$	$O(m)$
	Linear Regression (SGD)	$O(n_{epoch} nm)$	$O(m)$
	Logistic Regression (Binary)	$O(n_{epoch} nm)$	$O(m)$
	Logistic Regression (Multiclass OvR)	$O(n_{epoch} nmc)$	$O(mc)$
	Decision Tree	$O(n \cdot \log(n) \cdot m)$ $O(n^2 \cdot m)^*$ Worst case	$O(d_{tree})$
	Random Forest Classifier	$O(n_{trees} \cdot n \cdot \log(n) \cdot m)$	$O(n_{trees} \cdot d_{tree})$
	Support Vector Machines (SVMs)	$O(n^2 m + n^3)$	$O(m \cdot n_{SV})$
	k-Nearest Neighbors	—	$O(nm)$
$P(B A) = \frac{P(B \cap A)}{P(A)}$	Naive Bayes	$O(nm)$	$O(mc)$
	Principal Component Analysis (PCA)	$O(nm^2 + m^3)$	—
	t-SNE	$O(n^2 m)$	—
	KMeans Clustering	$O(iknm)$??

n: samples

m: dimensions

n_{epoch}: epochs

c: classes

d_{tree}: depth

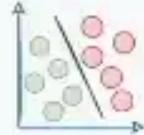
n_{SV}: Support vectors

k: clusters

i: iterations

Next...

5 Techniques to fine-tune LLMs



blog.DailyDoseofDS.com

LoRA-FA

$$h \in R^d$$

Frozen weights
Trainable weights

Pretrained weights
 $W \in R^{d*d}$

$$B \in R^{r*d}$$



$$A \in R^{d*r}$$

$$x \in R^d$$

$$h \in R^d$$



Pretrained weights

$$W \in R^{d*d}$$

$$B \in R^{r*d}$$



$$A \in R^{d*r}$$

$$x \in R^d$$

VeRA

$$h \in R^d$$

Frozen weights
Trainable vectors

Pretrained weights
 $W \in R^{d*d}$

$$b = 0$$

$$B \in R^{r*d}$$

$$d = 1$$

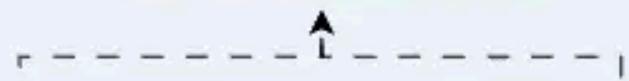
random and shared across layers

$$A \in R^{d*r}$$

$$x \in R^d$$

Delta-LoRA

$$h \in R^d$$



Pretrained weights

$$W \in R^{d*d}$$

$$W^{t+1} = W^t + c(A_{t+1} \cdot B_{t+1} - A_t \cdot B_t)$$

$$B \in R^{r*d}$$



$$A \in R^{d*r}$$

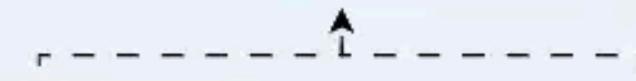
$$x \in R^d$$

Trainable weights

LoRA

Frozen weights
Trainable weights

$$h \in R^d$$



Pretrained weights

$$W \in R^{d*d}$$

$$B \in R^{r*d}$$



$$A \in R^{d*r}$$

$$x \in R^d$$

LoRA+

Almost similar to LoRA

LoRA update rule

$$A \leftarrow A - \alpha \frac{\delta J}{\delta A}$$

$$B \leftarrow B - \alpha \frac{\delta J}{\delta B}$$

LoRA+ update rule

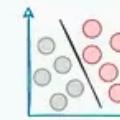
$$A \leftarrow A - \alpha \frac{\delta J}{\delta A}$$

Higher learning rate ~ for matrix B

$$B \leftarrow B - \lambda \alpha \frac{\delta J}{\delta B}$$

Next...

40 NumPy Methods That Data Scientists Use 95% of the Time



blog.DailyDoseofDS.com

NumPy Array Creation Methods

Method	Description
<code>np.array(<list>)</code>	NumPy array from Python list
<code>np.array(<list-of-lists>)</code>	NumPy array from list of lists
<code>np.array(<pandas-series>)</code>	NumPy array from PD Series
<code>df.values</code>	NumPy array from DataFrame
<code>np.zeros(<size>)</code>	NumPy array of all zeros
<code>np.ones(<size>)</code>	NumPy array of all ones
<code>np.eye(<size>)</code>	Identity NumPy array
<code>np.arange(<start>, <stop>, <step>)</code>	Equally spaced NumPy array with specific step
<code>np.linspace(<start>, <stop>, <count>)</code>	Equally spaced NumPy array with specific size
<code>np.random.randint(<low>, <high>, <size>)</code>	NumPy array of random ints
<code>np.random.random(<size>)</code>	NumPy array of random floats

NumPy Array Manipulation Methods

Method	Description
<code>array.reshape(<new-shape>)</code>	Reshape NumPy Array
<code>array.transpose() OR array.T</code>	Transpose NumPy Array
<code>np.concatenate(<np-arrays>, <axis>)</code>	Concatenate NumPy Arrays
<code>np.flatten(<Nd-np-array>)</code>	Flatten a NumPy Array
<code>np.unique(<np-array>, <axis>)</code>	Find unique elements
<code>array.tolist()</code>	NumPy Array to List

Search Methods

Method	Description
<code>np.argmax(<np-array>, <axis>)</code>	Max Element Index
<code>np.argmin(<np-array>, <axis>)</code>	Min Element Index
<code>np.where(<condition>, <true-return-value>, <false-return-value>)</code>	Conditional Search and Replacement
<code>np.nonzero(<np-array>)</code>	Index of non-zero elements

Mathematical Operations

Method	Description
<code>np.sin(<np-array>)</code>	
<code>np.cos(<np-array>)</code>	Trigonometric Functions
<code>np.tan(<np-array>)</code>	
<code>np.floor(<np-array>)</code>	Element-wise floor value
<code>np.ceil(<np-array>)</code>	Element-wise ceiling value
<code>np.rint(<np-array>)</code>	Round to nearest int
<code>np.round_(<np-array>, <decimal-places>)</code>	Round to decimal places
<code>np.exp(<np-array>)</code>	Element-wise exponent
<code>np.log(<np-array>)</code>	Element-wise logarithm
<code>np.sqrt(<np-array>)</code>	Element-wise square root
<code>np.sum(<np-array>, <axis>)</code>	Sum along an axis
<code>np.mean(<np-array>, <axis>)</code>	Mean along an axis
<code>np.std(<np-array>, <axis>)</code>	Std. dev along an axis

Matrix and Vector Operations

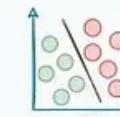
Method	Description
<code>np.dot(<np-array1>, <np-array2>)</code>	Dot Product
<code>np.matmul(<np-array1>, <np-array2>)</code>	
<code>np.array1 @ np.array2</code>	Matrix Multiplication
<code>np.linalg.norm(<np-array>)</code>	Vector Norm

Sorting Methods

Method	Description
<code>np.sort(<np-array>, <axis>)</code>	Sort Array
<code>np.argsort(<np-array>, <axis>)</code>	Return the order of indices that sort the array

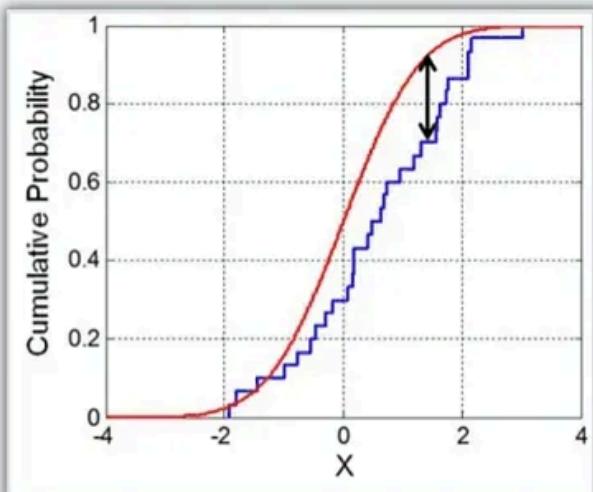
Next... 

11 Most Important Plots in Data Science

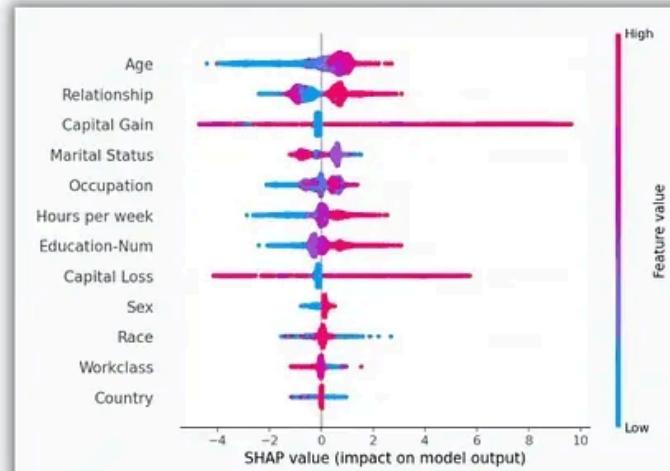


blog.DailyDoseofDS.com

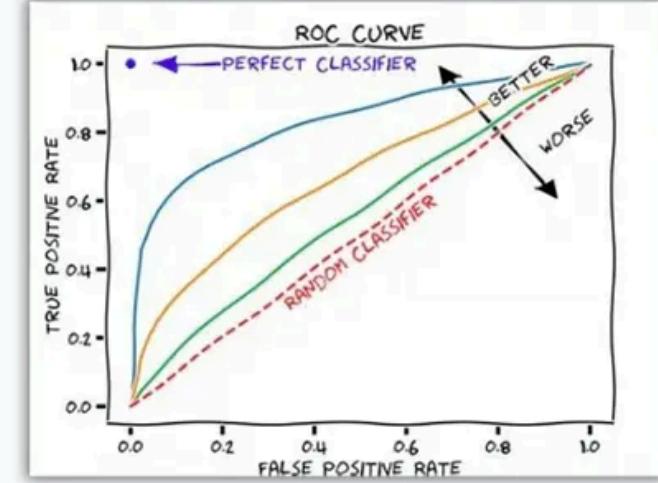
1 KS Plot



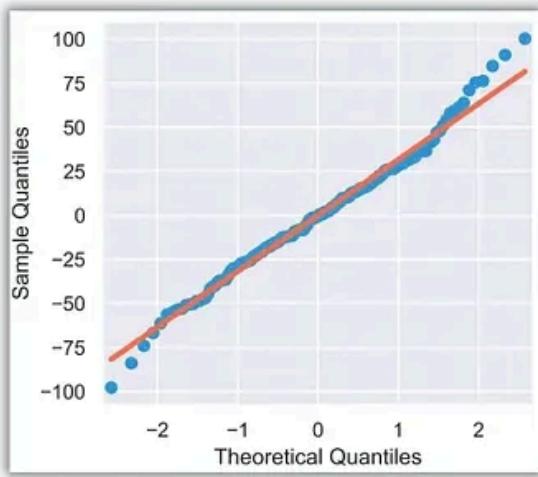
2 SHAP Plot



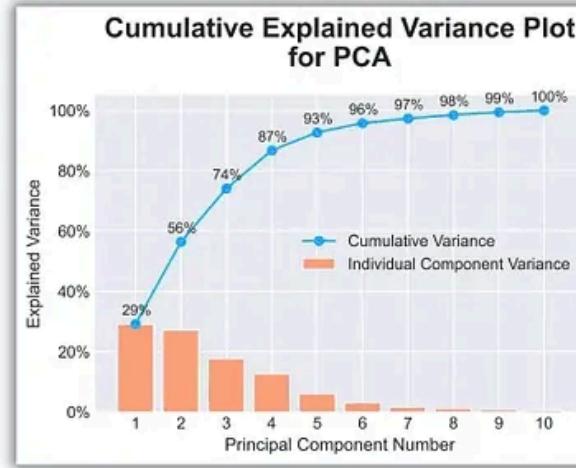
3 ROC Curve



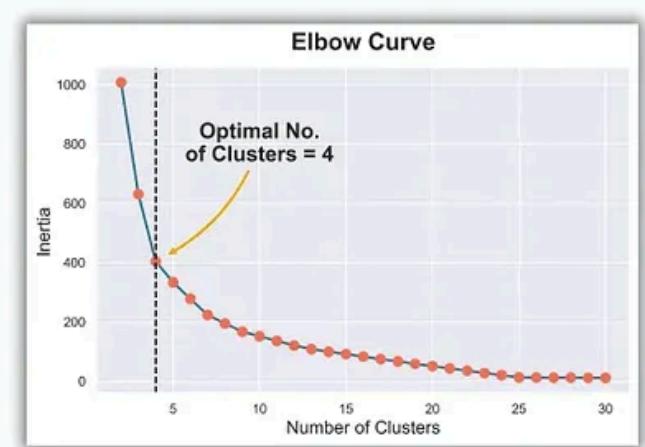
4 QQ Plot



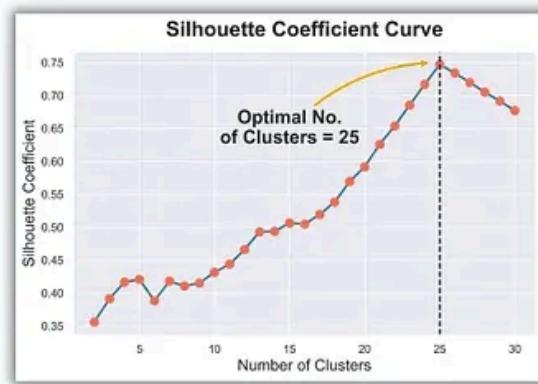
5 Cumulative Explained Variance



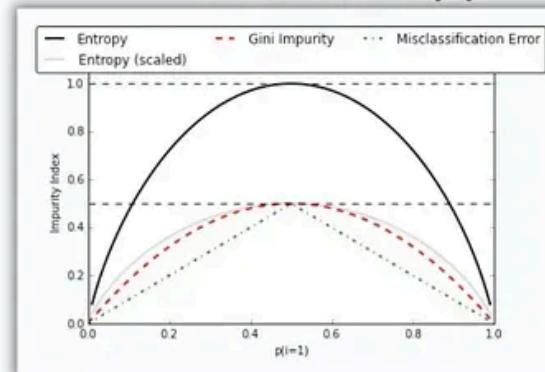
6 Elbow Curve



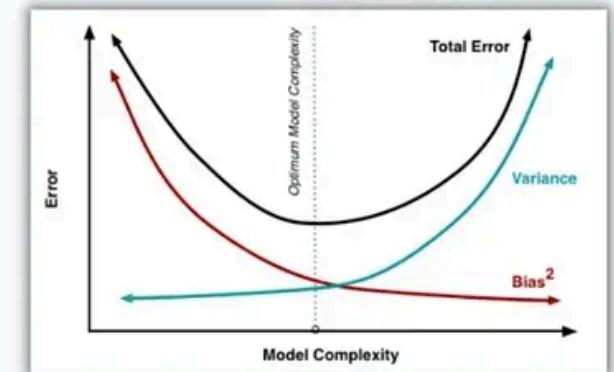
7 Silhouette Curve



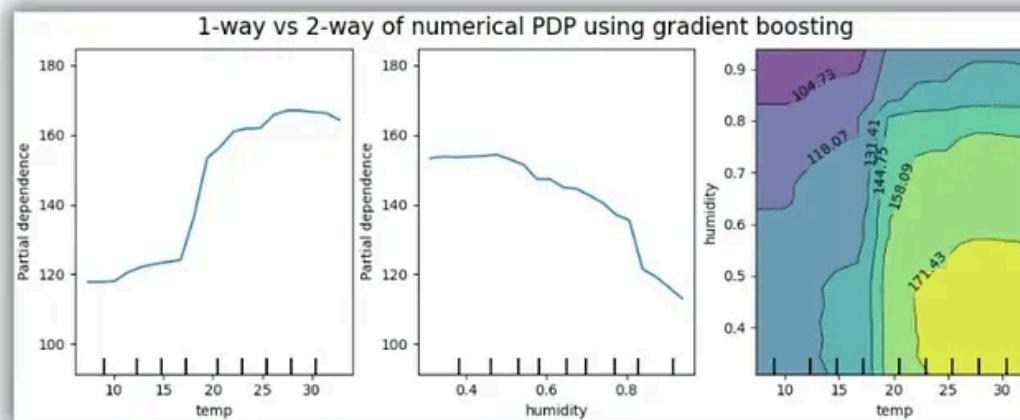
8 Gini impurity vs. Entropy



9 Bias-Variance Tradeoff



10 Partial Dependency Plot

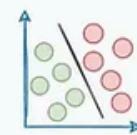


11 Precision Recall Plot



Next...

10 Most Common Loss Functions in Machine Learning



blog.DailyDoseofDS.com

Loss Function Name	Description	Function
--------------------	-------------	----------

Regression Loss Functions

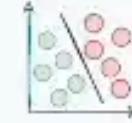
Mean Bias Error	Captures average bias in prediction. But is rarely used for training.	$\mathcal{L}_{MBE} = \frac{1}{N} \sum_{i=1}^N (y_i - f(x_i))$
Mean Absolute Error	Measures absolute average bias in prediction. Also called L1 Loss.	$\mathcal{L}_{MAE} = \frac{1}{N} \sum_{i=1}^N y_i - f(x_i) $
Mean Squared Error	Average squared distance between actual and predicted. Also called L2 Loss.	$\mathcal{L}_{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - f(x_i))^2$
Root Mean Squared Error	Square root of MSE. Loss and dependent variable have same units.	$\mathcal{L}_{RMSE} = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - f(x_i))^2}$
Huber Loss	A combination of MSE and MAE. It is parametric loss function.	$\mathcal{L}_{\text{Huberloss}} = \begin{cases} \frac{1}{2}(y_i - f(x_i))^2 & : y_i - f(x_i) \leq \delta \\ \delta(y_i - f(x_i) - \frac{1}{2}\delta) & : \text{otherwise} \end{cases}$
Log Cosh Loss	Similar to Huber Loss + non-parametric. But computationally expensive.	$\mathcal{L}_{\text{LogCosh}} = \frac{1}{N} \sum_{i=1}^N \log(\cosh(f(x_i) - y_i))$

Classification Loss Functions (Binary + Multi-class)

Binary Cross Entropy (BCE)	Loss function for binary classification tasks.	$\mathcal{L}_{BCE} = \frac{1}{N} \sum_{i=1}^N y_i \cdot \log(p(x_i)) + (1 - y_i) \cdot \log(1 - p(x_i))$
Hinge Loss	Penalizes wrong and right (but less confident) predictions. Commonly used in SVMs.	$\mathcal{L}_{\text{Hinge}} = \max(0, 1 - (f(x) \cdot y))$
Cross Entropy Loss	Extension of BCE loss to multi-class classification.	$\mathcal{L}_{ce} = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^M y_{ij} \cdot \log(f(x_{ij}))$ N : samples; M : classes
KL Divergence	Minimizes the divergence between predicted and true probability distribution	$\mathcal{L}_{KL} = \sum_{i=1}^N y_i \cdot \log\left(\frac{y_i}{f(x_i)}\right)$

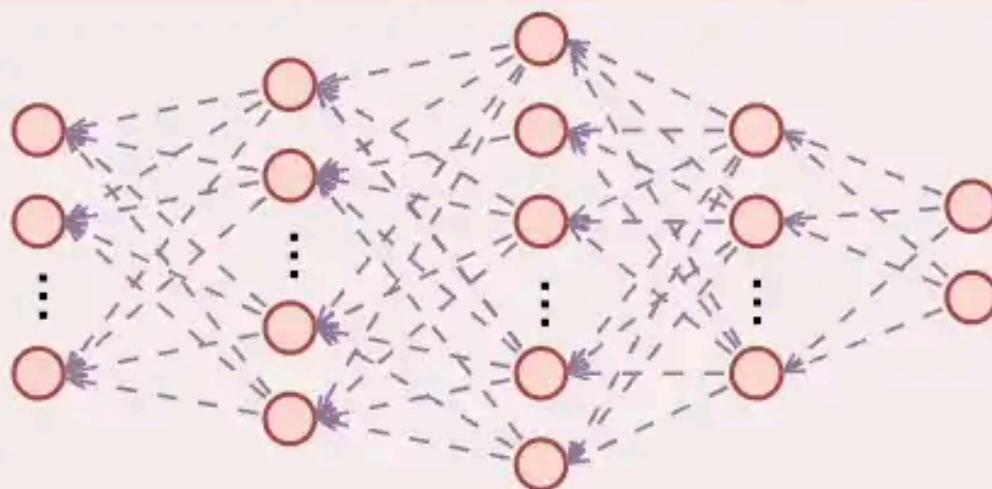
Next... 

Full Fine Tuning, LoRA and RAG

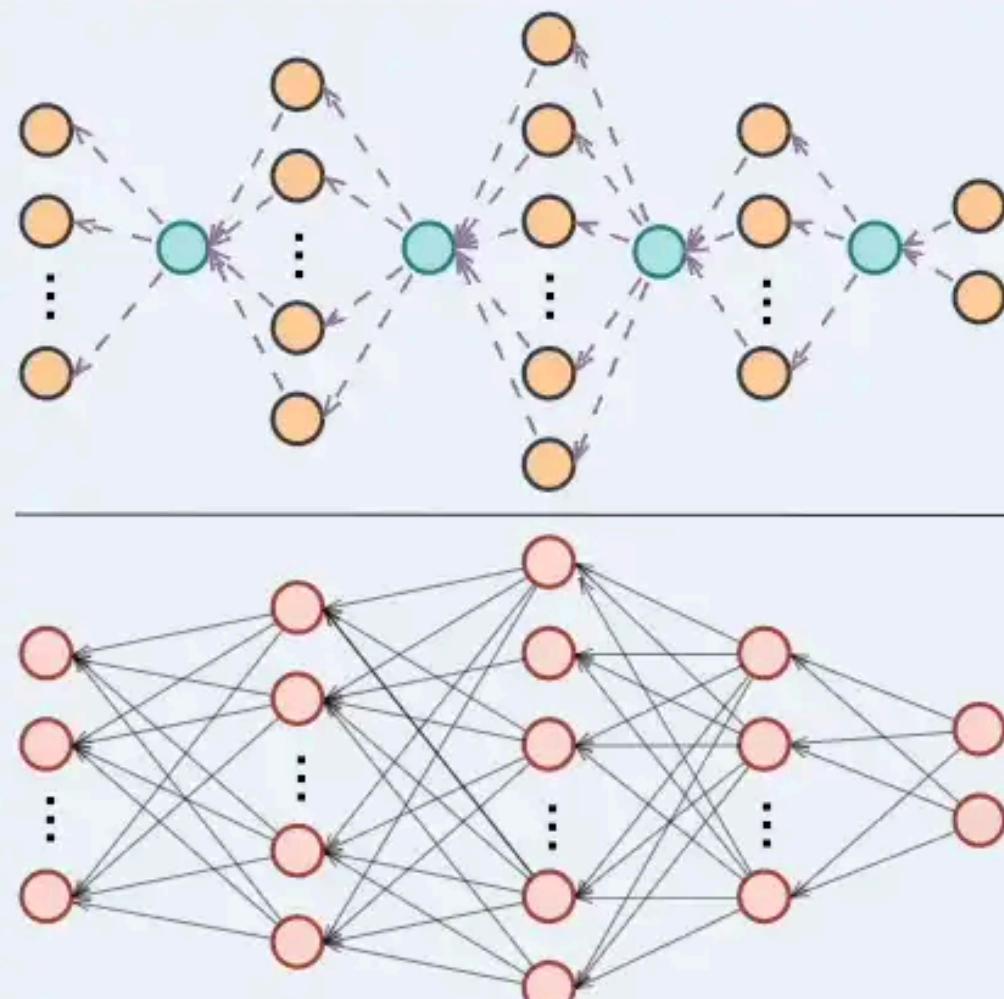


blog.DailyDoseofDS.com

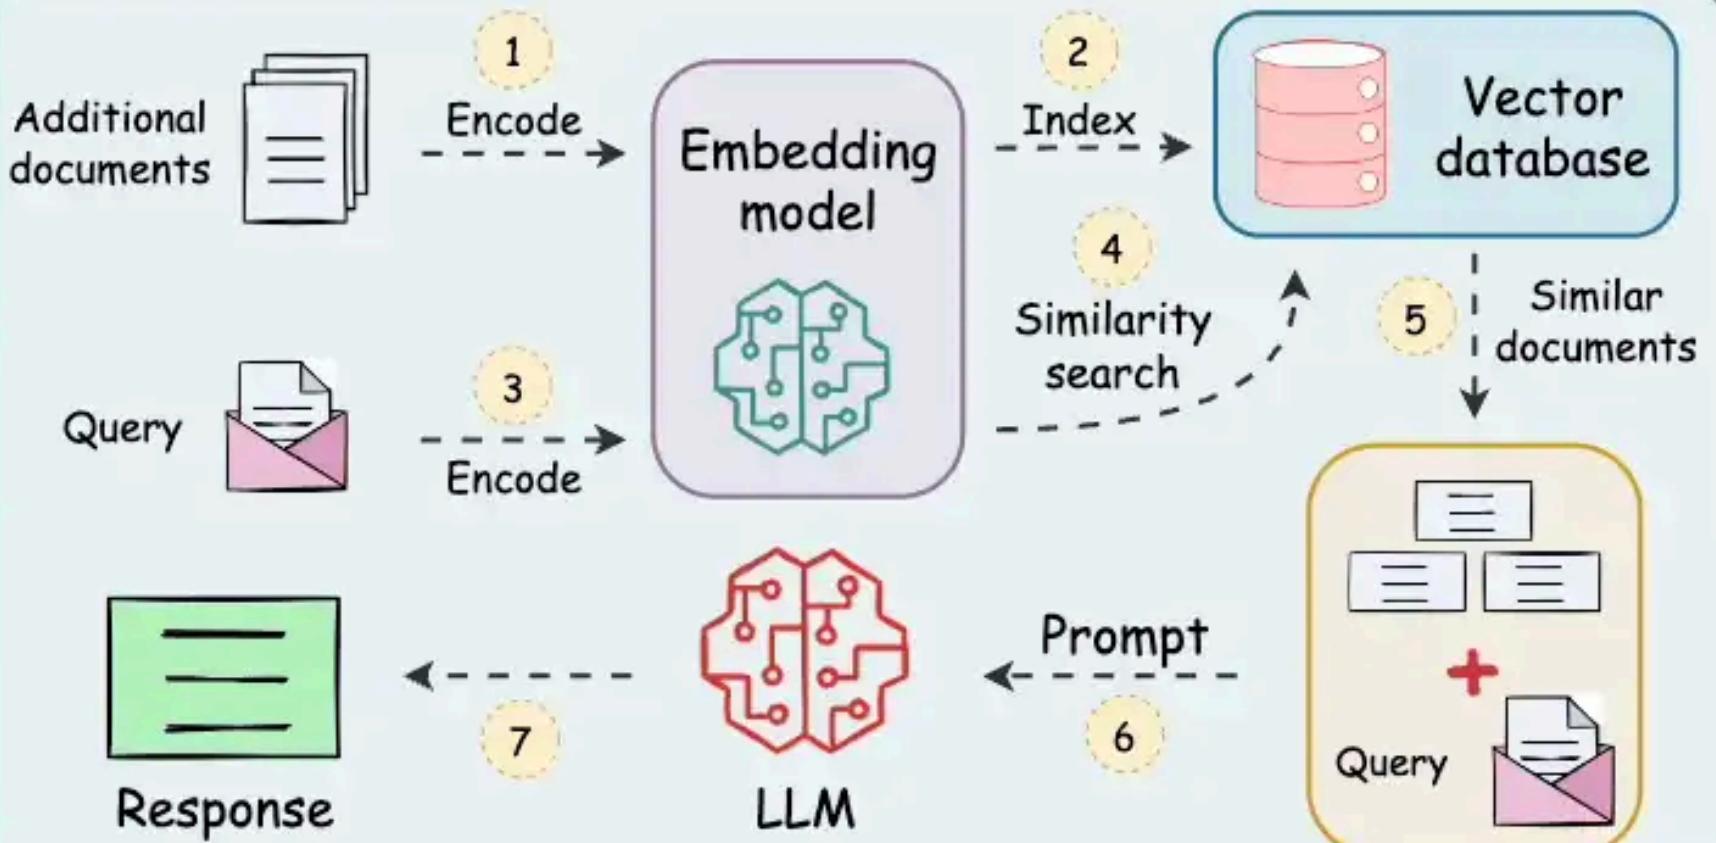
Full Fine Tuning



LoRA Fine Tuning

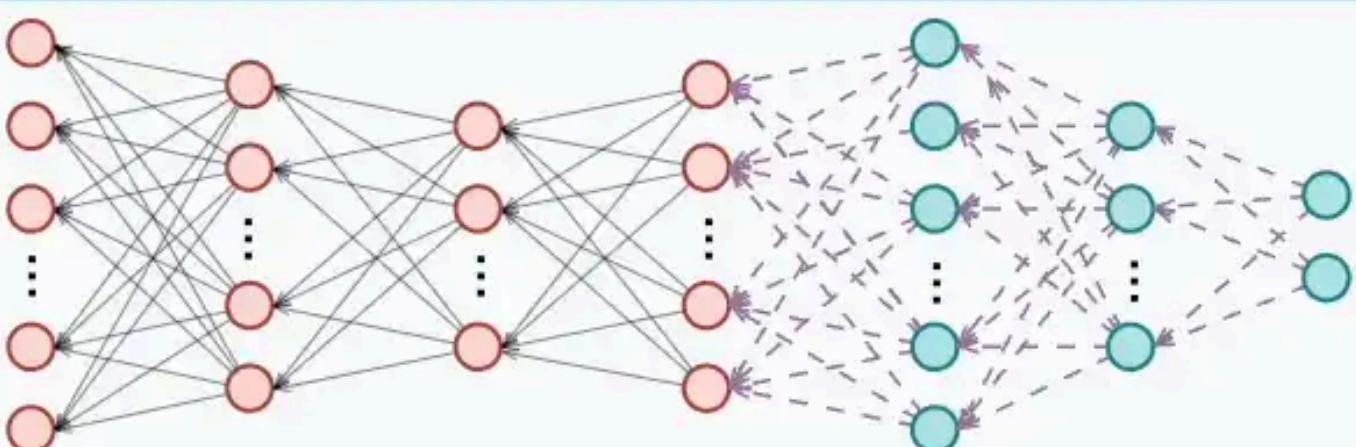


RAG



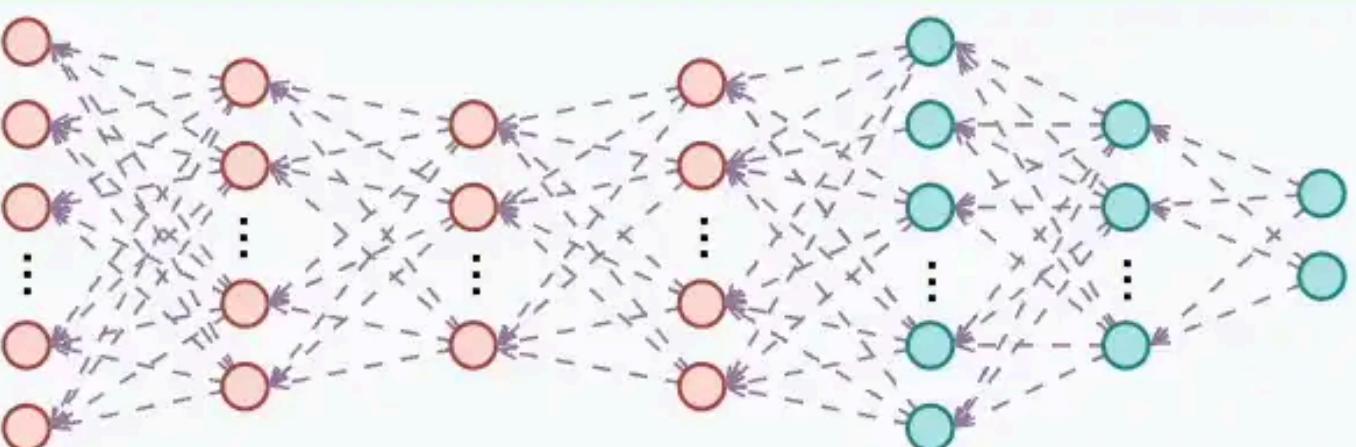
Next...

Transfer Learning



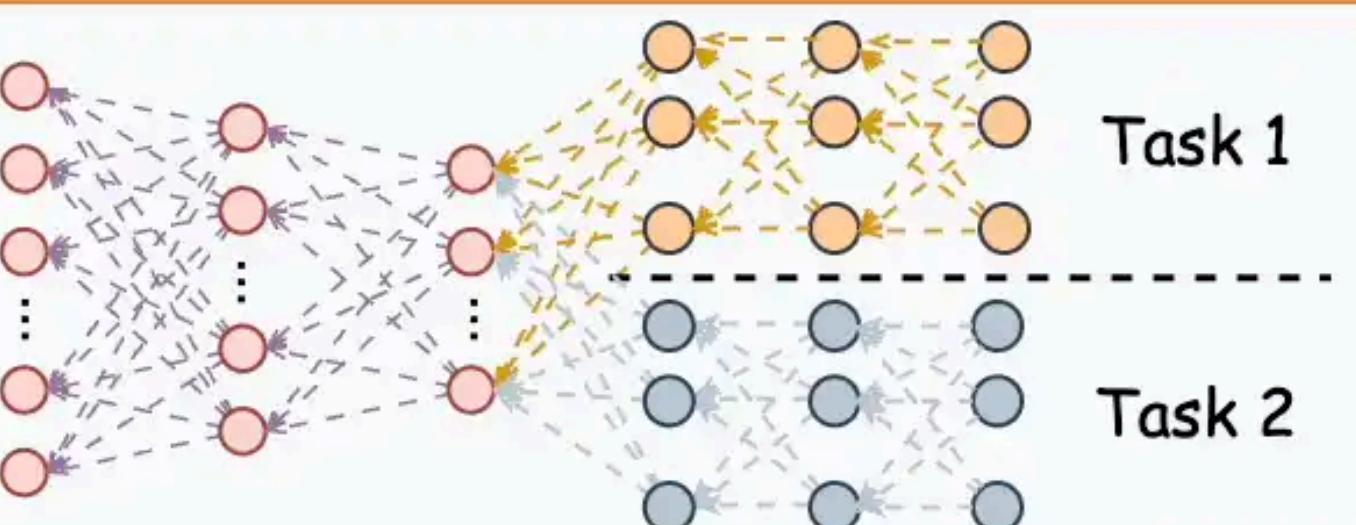
- ↔--- Gradient flow
- ← No gradient flow
- Neurons from a pre-trained model
- Appended neurons

Fine Tuning



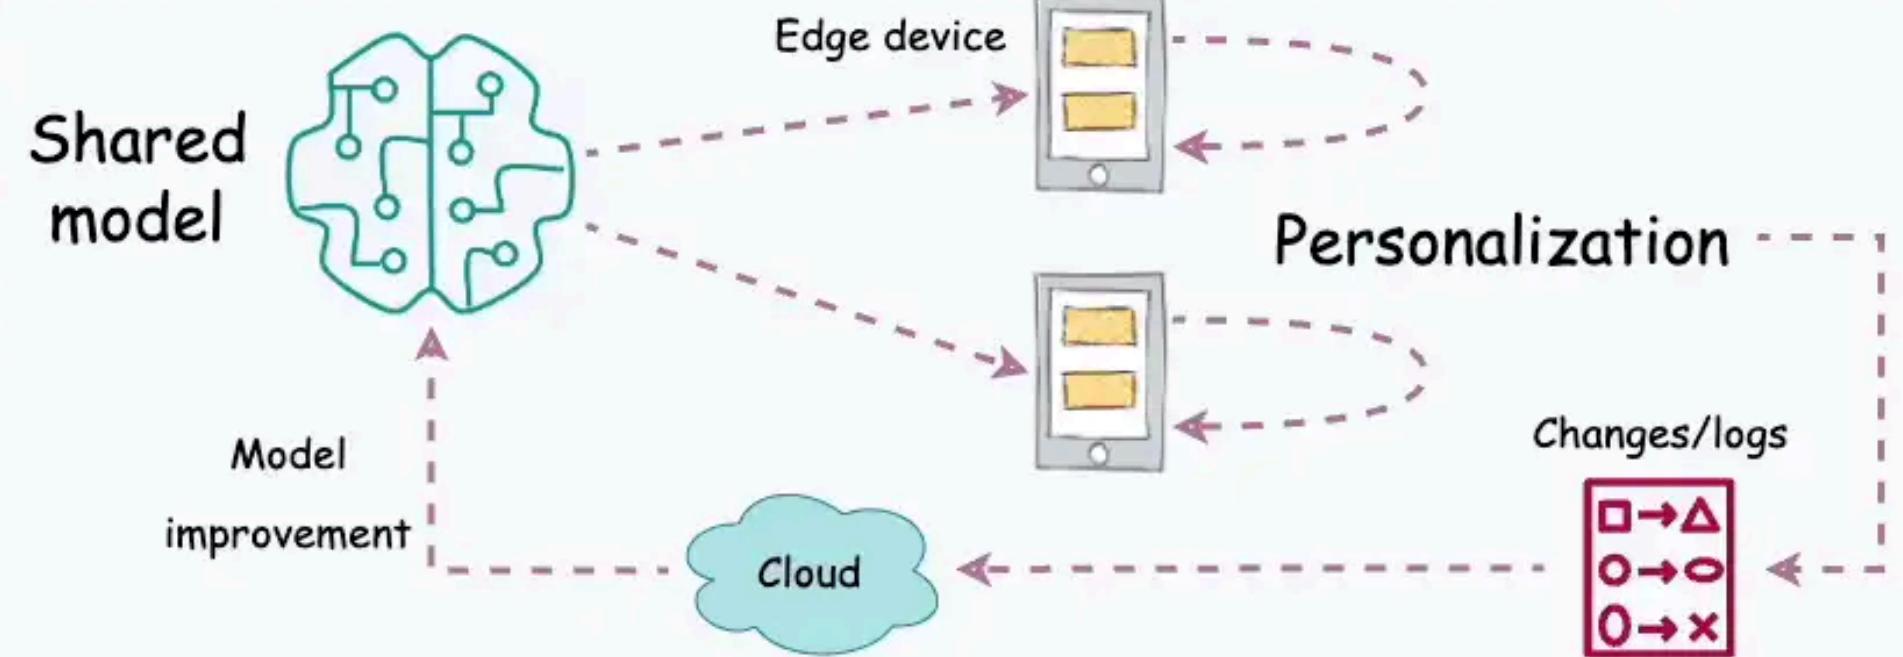
- ↔--- Gradient flow
- Full pre-trained network
- Appended network

Multitask Learning



- ↔--- Gradient flow
- Shared network
- Task-specific branches

Federated Learning



Next...